



ARCHITECTURE VISION

eDiscovery and Litigation Document Review Platform

Version 0.2

For

International Law Firm

Contents

1 Introduction	4
1.1 Purpose	4
1.2 Scope	4
1.3 Definitions	5
1.4 References	6
1.5 Documentation Roadmap	6
2 Executive Summary	8
2.1 Key Decisions	8
2.2 Key Risks and Open Issues	9
3 Architectural Drivers	10
3.1 Business Case	10
3.1.1 Business Goals	11
3.1.2 Major Features	12
3.2 Use Case Model	13
3.2.1 Use Case View <View Name>	13
3.3 Design Constraints	15
3.4 Quality Attribute Scenarios	15
4 Solution Architecture	18
4.1 Big Picture	18
4.1.1 Solution Context	18
4.1.2 Solution Decomposition	19
4.2 Design Concerns	21
4.3 Development Technology Stack	21
4.3.1 Development Languages, Frameworks, and Libraries	21
4.3.2 Development Tools	22
4.3.3 External Integration Points	22
5 Implementation Roadmap	24
5.1 Implementation Deliverables	24
5.2 Implementation Milestones	24
5.3 Estimate	25
5.3.1 Assumptions and Limitations	25
5.3.2 Estimate	25
5.4 Team	26
5.4.1 Team Skillset	26
5.4.2 Team Structure	27

Revision History

Date	Version	Description	Author
10/09/2026	0.1	Initial draft	Maksym Halych
10/09/2026	0.2	Added key decisions, key risks and the business case	Maksym Halych

1 Introduction

1.1 Purpose

This Architecture Vision elicits the significant architecture drivers such as business, functional, nonfunctional requirements and constraints, defines architecture, and develops a roadmap for the eDiscovery and Litigation Document Review Platform implementation. The document is intended as a primary technical guidance into solution implementation for the development team.

The document answers the Client's request for proposal, listed in References. The Client is an international law firm with 1,400 lawyers in 14 offices, whose litigation, arbitration and regulatory-investigation practices process roughly 220 million documents per year across 300 active matters. Document review is outsourced today to three external hosting vendors at a cost of EUR 18.5 million per year, and the Client intends to replace them with a firm-operated eDiscovery platform hosted under its own security controls.

The solution architecture is designed following the guiding principles outlined in Appendix A — Architecture Design Methodology.

1.2 Scope

The document describes the proposed architecture of the eDiscovery and Litigation Document Review Platform towards development of the solution that will satisfy business, functional, non-functional requirements and constraints provided by the Client. This Architecture Vision covers the following information:

- Significant architectural drivers for the Solution
- Proposed software architecture of the solution based on these drivers
- Technology stack and environment definitions
- Operations specific perspectives
- Development road map and high level estimates for effort, team size and skill sets.

The solution covers the eDiscovery lifecycle from legal hold to production, in the three delivery phases the Client requested:

- Phase 1: legal hold, collection connectors, processing and early case assessment
- Phase 2: the review workspace with continuous active learning, privilege analytics, redaction and production
- Phase 3: client reporting portals and advanced analytics such as communication mapping

The scope also includes the migration of ten historical matters from the external hosting

vendors, and the hosting of each matter in the region its client approves: the EU, the UK or the US.

1.3 Definitions

The Definitions section lists the acronyms and terms used in this document which might possess lesser familiarity or double meaning to the reader.

#	Term	Definition
1.	Continuous active learning	A machine-learning review method that retrains the model on every reviewer decision and keeps the likely relevant documents first in the queue.
2.	Custodian	A person who keeps data that is related to a matter, for example an employee with a related mailbox.
3.	Early case assessment	Analysis of collected data before review to estimate its volume, cost and key facts and to cull irrelevant data.
4.	eDiscovery	Electronic discovery: the identification, preservation, collection, processing, review and production of electronically stored information for litigation or investigations.
5.	Ethical wall	An access barrier. It keeps persons with a conflict of interest away from the data of a matter.
6.	Go-live	The first production release of the platform with real client matters. Some objectives start their time frame at go-live.
7.	Legal hold	A notice that instructs the persons who hold data relevant to a matter to preserve it, and suspends its routine deletion.
8.	Matter	A single litigation, arbitration or regulatory investigation the firm handles for a client; the unit of data, access and billing.
9.	Precision	The percentage of the documents that the model marks as relevant and that are really relevant.
10.	Privilege	Legal professional privilege: protection that allows a document, such as lawyer-client advice, to be withheld from production.
11.	Production	Delivery of relevant, non-privileged documents to the opposing party or the court in an agreed format; distinct from the production environment the platform runs in.
12.	Recall	The percentage of all relevant documents in a matter that the model finds.

#	Term	Definition
13.	Redaction	Masking privileged or sensitive content in a document before production.
14.	TAR	Technology-assisted review: a review method in which a machine-learning model learns from reviewer decisions and sorts documents by relevance.

1.4 References

The References section provides a complete list of all the documents referenced elsewhere in this document.

#	Version	Date	Document Name	Published by
1.	1.0	—	Request for Proposal: eDiscovery and Litigation Document Review Platform	International Law Firm (the Client)

1.5 Documentation Roadmap

This section guides into the Architecture Vision document structure to help find the information of interest.

- Executive Summary speaks about the proposed architecture on the highest level and covers:
 - Key Decisions about the architectural and operational choices made for the solution
 - Key Risks and Open Issues detected with the analysis of the available context, requirements, and proposed decisions
- Architectural Drivers elicits the known essential requirements and constraints for the solution to design which play key role in forming architectural decisions and overall architecture.
 - Business Case describes the solution from the business standpoint including business goals and objectives, and major features
 - Service Level Agreement defines the important technical constraints and guarantees under which the solution will be serviced to its clients
 - Use Case Model lists the key Use Cases for the designed solution
 - Domain Model shows the key business entities with attributes and relationships between them.

- Design Constraints include business, resource, technical and other constraints accounted for in the architecture of the solution.
- Quality Attribute Scenarios are a set of the testable scenarios clarifying non-functional requirements for the system quality attributes such as performance, maintainability, and others.
- Solution Architecture defines the proposed architecture as a set of architectural views in the format defined in Appendix C — How View is Documented.
 - Big Picture shows the solution architectural context, high level decomposition into components, and followed reference architecture.
 - Development Technology Stack selects the tools, frameworks, libraries, external services and other technologies the solution implementation will rely on.
- Implementation Roadmap proposes the plan for solution implementation including:
 - Implementation Deliverables expected to be implemented and delivered at the implementation phase.
 - Implementation Milestones tied to the project timeline
 - Estimate of complexity/size, effort, schedules for the implementation
 - Team skillset and structure based on the technology, competence, and schedule requirements

2 Executive Summary

The section Executive Summary highlights key architectural decisions made for the solution described in [Business Case](#). These decisions are defined and discussed in depth in [Solution Architecture](#) while [Implementation Roadmap](#) lays out the proposed milestones, estimates, and team to implement the provided architecture vision.

This section also summarizes the key business and technical risks related to the solution implementation. These risks are uncovered in depth in the rest of the document.

2.1 Key Decisions

The section outlines key design decisions about the solution including the architecture big picture and most essential technologies and external services to rely on.

The solution will be a firm-operated eDiscovery platform on Amazon Web Services (AWS). The key decisions are:

- **Cloud provider.** The platform uses AWS, because the delivery team has strong AWS experience. This decreases the delivery risk in the 16-month program. The platform connects to Microsoft 365 and Purview through the Microsoft Graph API. The Client's Microsoft Entra ID gives single sign-on (SSO) with multi-factor authentication (MFA).
- **Data residency.** Each region (EU, UK and US) has an independent platform instance in its own AWS accounts. Service control policies block all other regions. Thus, the data of a matter stays in the region that its client approves. A separate AWS KMS key encrypts the data of each client.
- **Build approach.** The team develops the eDiscovery workflow. AWS managed services and open-source libraries do the generic functions: Apache Tika extracts text and metadata, Amazon Textract does optical character recognition (OCR), and Amazon OpenSearch Service does the full-text search. Thus, there are no license costs and no new vendor dependency. Sources without a standard connector, such as mobile devices, come in as exports from forensic tools.
- **Processing and storage.** Amazon SQS, AWS Step Functions and Amazon ECS with AWS Fargate process the collected data. The number of workers increases to process 2 TB of raw data each day. Amazon S3 keeps the documents, and Amazon Aurora PostgreSQL keeps the matter data.
- **Review.** Technology-assisted review (TAR) uses continuous active learning and statistical validation on samples. Courts in the UK, the US and some EU jurisdictions accept this method. Large language models in Amazon Bedrock help with privilege detection, through endpoints in the same region only. A reviewer always makes the final privilege decision.
- **Audit.** An append-only, hash-chained ledger records all user and system actions. Copies

of its digests go to S3 Object Lock storage. Thus, any change to the audit trail is visible.

2.2 Key Risks and Open Issues

The section lists the key risks related to the solution design and implementation. It also lists key open issues where architectural decisions have not been made yet or are likely to change.

#	Risk Description	Mitigation Strategy
R-1	Custodian data is usually in the Microsoft 365 tenants of the clients, not in the firm's tenant.	Confirm the data sources in Phase 1. Design connectors for many client tenants.
R-2	Open-source libraries possibly cannot read some rare formats, such as Lotus Notes archives and phone images.	Import these sources as exports from forensic tools. Record the hash value and metadata of each item.
R-3	TAR can miss its recall and precision targets. Court acceptance of TAR is different in each jurisdiction.	Keep linear review available. Agree on the validation protocol with the court.
R-4	Privilege detection can miss privileged documents. These documents then go to the opposing party.	A reviewer makes the final decision. Review a sample of the unflagged documents before production.
R-5	The limits of 16 months, EUR 6 million and 12 people are tight for 3 phases and 3 regions.	Deploy all regions from one infrastructure-as-code template. Give priority to Phases 1 and 2.
R-6	Model validation is possible only with real matters from 3 large clients in the first year.	Get the agreement of these clients in Phase 1. Use the migrated matters as more validation data.
R-7	Client data cannot go on the platform before the security assessment passes. A late failure moves go-live.	Design the security controls from the start. Do a pre-assessment before the formal assessment.
R-8	If go-live is late, the vendor contracts end before the platform is ready.	Align go-live with the 2027 contract renewals. Negotiate short extensions as a fallback.

3 Architectural Drivers

The section captures significant requirements driving the solution architecture and road map. The requirements which are not influencing the solution architecture in major ways and low level requirement details and scenarios are typically excluded from this section and can be found in the requirement specification or the product backlogs.

3.1 Business Case

The section lays out the business case for the eDiscovery and Litigation Document Review Platform.




Figure 3.1. Business level view of the eDiscovery and Litigation Document Review Platform

Figure 3.1 shows the users of the platform, its external systems, and the data flows between them.

For its matters, the Client must preserve, collect, review and produce electronically stored information. Today, three external vendors do this work with different workflows, and contract reviewers read every document. This costs EUR 18.5 million each year. For 68 percent of matters, the review cost is more than the budget. Partners cannot compare cost and progress across matters.

The platform puts the full eDiscovery lifecycle into one firm-operated system:

- The matter-intake and conflicts system sends each new matter and its ethical walls to the platform.

- eDiscovery analysts send legal holds to custodians. Custodians acknowledge the holds on the platform.
- The platform preserves and collects custodian data from corporate and cloud sources, such as Microsoft 365 and Purview.
- The platform processes the data and removes duplicates. Continuous active learning shows the most relevant documents first to review managers and contract reviewers.
- Privilege analytics find privileged content for redaction. Then the platform produces the documents to opposing parties and courts.
- Partners monitor the cost and progress of each matter in a browser or on a tablet. Cost data goes to the practice-management and billing system.
- The platform exchanges documents with the iManage Work document management system. Security administrators manage access and ethical walls.
- In Phase 3, client counsel monitor their matters in a client reporting portal.

Each matter stays in the region that its client approves: the EU, the UK or the US. Per-client keys encrypt the data, and an audit trail records all actions. Thus, the firm can show clients and courts a secure, defensible and auditable review process. The firm can also offer fixed fees, because it controls its review costs.

3.1.1 Business Goals

The section enumerates essential business goals for the solution.

#	Description
BG-1	Decrease the document review cost and keep the review margin in the firm
BG-2	Offer fixed-fee litigation services, to keep and win panel appointments
BG-3	Make the review process defensible in court, from legal hold to production
BG-4	Keep client data secure and in the jurisdictions that the clients approve
BG-5	Stop the dependence on external hosting vendors

A business goal is not measurable. Thus, the table below gives measurable objectives for the goals. One goal can have many objectives, and one objective can help many goals. Each time frame starts at program start or at go-live.

#	Description	Business Goals
BO-1	Decrease the review cost per collected document by 45 percent, from approximately EUR 0.084, in 24 months after go-live. At today's volume, this is approximately EUR 8.3 million each year	BG-1, BG-2
BO-2	Decrease human review from 100 to 40 percent of documents on matters above 500,000 documents, in 24 months after go-live	BG-1, BG-2
BO-3	Get a TAR recall of 75 percent or more at a precision of 65 percent or more on validation samples. Use a statistical method that courts accept	BG-3
BO-4	Decrease the time to issue a legal hold from 5 business days to less than 4 hours. Track 100 percent of custodian acknowledgements. Do both in 18 months from program start	BG-3
BO-5	Find 95 percent or more of the privileged documents in sample sets. Thus, no privileged document goes to the opposing party	BG-3, BG-4
BO-6	Pass an independent ISO 27001-aligned security assessment and the security questionnaires of the 10 largest clients. Do this before the platform keeps client data	BG-4
BO-7	Put 90 percent of new matters on the platform by month 20 from program start	BG-5

3.1.2 Major Features

The section enumerates solution major features.

#	Description
F-1	Legal hold: send hold notices to custodians, track acknowledgements, and preserve data in Microsoft 365 and Purview (Phase 1)
F-2	Collection: standard connectors for corporate and cloud sources, and a bulk import API for other sources (Phase 1)

#	Description
F-3	Processing and early case assessment: text and metadata extraction, OCR, de-duplication, and volume and cost estimates (Phase 1)
F-4	Review workspace with TAR and continuous active learning, in a browser and on a tablet (Phase 2)
F-5	Privilege detection and redaction (Phase 2)
F-6	Production of documents to opposing parties and courts in court-ready formats (Phase 2)
F-7	Matter security: SSO with MFA, matter-level ethical walls, per-client keys, regional hosting, audit trail and defensible deletion (all phases)
F-8	Dashboards for cost, progress, capacity and review throughput, and cost export to the billing system (Phases 1 and 2)
F-9	Client reporting portal, and communication maps between custodians (Phase 3)
F-10	Migration of 10 historical matters from the external vendors (Phase 2)

3.2 Use Case Model

The section lists the key Use Cases for the designed solution as one or more Use Case views.

3.2.1 Use Case View <View Name>

3.2.1.1 View Context

The interaction outlined in the feature [F-3](#) happens between the user and the user facing we application

3.2.1.2 Representation

(Add your Use Case diagrams here. They should cover all major cases of the system usage (at least 5 use cases). You could have a few separate diagrams to cover all the use cases)




Figure 3.2. <Use Case View Name>

(add description for your diagrams here, refer to features using IDs you defined in 3.1.2.)

This view defines primary Use Cases for the feature F-3, corresponding to the types of interaction between the user and the web application supported by the solution.

3.2.1.3 Element Catalog

(Describe your actors and use cases in the table below)

#	Name	Description
ACT-1	User	A business user of the system
UC-1	Use Case1	The user logs into the system.

#	Name	Description
UC-2	Use Case2	The user pays with the registered credit card.
<id>	<element name>	<Use Case description or scenario>

3.3 Design Constraints

The section lists the constraints accounted for in the designed solution. These can be of business, technical, resource, and other types.

(list identified constraints in the table below)

#	Description
CONST-1	A minimum of 1000 simultaneous users must be supported
CONST-2	Time to market must be within one year
CONST-3	Java+RDBMS on Amazon technology stack must be used
CONST-4	The existing client's dev team skill set is based on Spring framework and JPA
<constraint id>	<constraint description>

3.4 Quality Attribute Scenarios

A Quality Attribute Scenario is an unambiguous and testable requirement for one or more Solution Quality Attributes such as Performance, Usability, Maintainability and others. The scenario consists of six parts: Source of Stimulus, Stimulus, Environment, Artifact, Response, testable and accurate Response Measure.

This section lists and prioritizes the scenarios pertinent to the designed solution.

(List your quality attributes in the table below (at least 5))

#	Quality Attribute	Scenario	Business Priority	Related To
QA-1	Security (credentials transport security)	At all times the credentials entered by the user during log-in are transferred to the server over encrypted, secure channel without the chance of sniffing by third party.	High	<u>UC-1</u>
QA-2	Usability (easiness of payment)	When logged in and navigated to the payment page it takes the user up to 3 clicks to pay with the preregistered valid credit card	Medium	<u>UC-2</u>
<scenario id>	<attribute name (scenario meaning)>	<quality attribute scenario description>	<priority level: High, Medium, Low>	<use case, feature, constraint ids>

(At least one QA from the table above should be documented in the graphic format as below)

Scenario: QA-1

Quality Attributes: Security

Business Priority: High

Related To: UC-1

Description: At all times the credentials entered by the user during log-in are transferred to the server over encrypted, secure channel without the chance of sniffing by third party.

Environment: Normal operation conditions




Figure 3.3. Quality attribute scenario QA-1 in graphic form

4 Solution Architecture

The section Solution Architecture is primary for the Architecture Vision document. It defines and reasons about the solution architecture design based on requirements and constraints identified in the section Architectural Drivers.

4.1 Big Picture

The section includes a list of architectural views covering the designed solution along with the context it runs in on the high level.

4.1.1 Solution Context

4.1.1.1 Intent

The view defines the primary solution components collaborating with the external systems and services. It is driven by the Business Case.

4.1.1.2 Representation

(add your C4 context diagram below)




Figure 4.1. Solution Context diagram

(explain your context diagram)

4.1.1.3 Element Catalog

Table of annotated elements.

(annotate elements from your diagram in the table below)

Name	Description
Cloud Based Solution	Responsible for implementing the REST API to serve data on request from the mobile and web clients
Element2	Responsible for a, b, c
Element3	Responsible for a, b, c
<element name>	<element description and responsibilities>

4.1.2 Solution Decomposition

4.1.2.1 Intent

The view defines the runtime decomposition of the server-side part of the solution. It is driven by the Business Case and the architecture best practices applicable to the cloud-based applications.

4.1.2.2 Context

The view context is defined by the view Solution Context where this section represents the decomposition of the system using the Containers view.

4.1.2.3 Representation

(add your container diagram below)




Figure 4.2. Cloud Solution Container view

(explain your container diagram)

The cloud based part of the solution is decomposed into several parts documented below combining several standard architectural patterns applicable to the highly loaded cloud based SaaS applications: Load Balancer, Data Cache, Background Process, Shared Storage, and Static Content Provider. The subsection reasoning provides detailed discussion of these choices.

4.1.2.4 Element Catalog

Table of annotated elements.

(annotate elements from the container diagram here)

#	Description
Element1	Responsible for a, b, c
Element2	Responsible for a, b, c
Element3	Responsible for a, b, c

#	Description
<element name>	<element description and responsibilities>

4.2 Design Concerns

The section lists the concerns defined for the solution.

(list your design concerns/ architecturally significant requirements in the table below)

#	Description
CONC-1	<concern description>
CONC-2	<concern description>
CONC-3	<concern description>
CONC-4	<concern description>
<concern id>	<concern description>

4.3 Development Technology Stack

The section includes a list of decisions on tech stack, frameworks, libs, tools, external services, etc.

4.3.1 Development Languages, Frameworks, and Libraries

4.3.1.1 Intent

The view lists the set of programming languages, frameworks, and libraries the solution implementation will depend on.

4.3.1.2 Context

The context is provided by the view [Solution Context](#).

4.3.1.3 Element Catalog

Table of annotated elements.

(list all technologies, frameworks, libraries that have to be used in the solution in the table below)

Name	Version	Description
Framework1	x.x	Responsible for a, b, c

Name	Version	Description
Library2	3.0-RC1	Responsible for a, b, c
Library3	x.x.x.x	Responsible for a, b, c
<name>	<version>	<description and responsibilities>

4.3.2 Development Tools

4.3.2.1 Intent

The view lists the set of tools the development team will rely upon in solution implementation.

4.3.2.2 Context

The context is provided by the view Solution Context.

4.3.2.3 Element Catalog

Table of annotated elements.

(list all the tools, like IDEs, SAAS, plugins, etc that will be used during the development)

Name	Version	Description
Framework1	x.x	Responsible for a, b, c
Library2	3.0-RC1	Responsible for a, b, c
Library3	x.x.x.x	Responsible for a, b, c
<name>	<version>	<description and responsibilities>

4.3.3 External Integration Points

4.3.3.1 Intent

The view lists the external systems and services the solution implementation will integrate with.

4.3.3.2 Context

The context is provided by the view Solution Context.

4.3.3.3 Element Catalog

(list all the integration points if you have any 3-rd party integration in the table below)

Name	Version	Description
Framework1	x.x	Responsible for a, b, c
Library2	3.0-RC1	Responsible for a, b, c
Library3	x.x.x.x	Responsible for a, b, c
<name>	<version>	<description and responsibilities>

5 Implementation Roadmap

Implementation Roadmap defines the solution implementation road map including list of implementation deliverables, major milestones, effort estimates, and recommended team skillset, structure, and size.

5.1 Implementation Deliverables

The implemented solution will include the following parts as deliverables:

(list your deliverables in the table below)

Name	Refer to Requirements	Refer to Design
Application 1		
Service 2		
Mobile Client 3		
Production Deployment 4		

5.2 Implementation Milestones

The section proposes the major milestones to guide the solution implementation.

(decide and list main milestones for your project in the table below)

Milestone	Description	Outcomes
M0	Project bootstrap	• Dev Team ramped up • Dev and Testing Environment set up • Project Skeleton built
M1	Main application with web front-end development	• Alpha version of the web application • Deployment to the staging

Milestone	Description	Outcomes
M2	Development and stabilization of the web application, development of the deployment and monitoring procedures, release	• Release 1.0 of the application prepared and stable • DevOps framework functional • Release 1.0 goes to production

5.3 Estimate

The Estimate section provides the estimates for the solution implementation based on the proposed architecture and selected estimation methodology.

5.3.1 Assumptions and Limitations

This section describes the known limitations imposed by selected platform, technology, hardware, operating system, third party components, etc. which might affect implementation effort and schedule estimates.

(Describe assumptions and possible limitations that could affect delivery timelines in the table below)

#	Description	Responsible
A1	Assumption 1	<Client Name>
A2	Assumption 2	SoftServe Team
L1	Limitation 1	<Third Party Name>

5.3.2 Estimate

The Estimate section provides effort and, optionally, schedule estimates for the implementation based on the architecture proposed in this Vision document using the WBS approach.

(list delivery items and estimation for these items to be ready)

#	Item	Complexity	Min Effort	Max Effort
1.				

#	Item	Complexity	Min Effort	Max Effort
2.				
3.				
4.				
5.				
Total:				

5.4 Team

The Team section outlines the proposed team skillset and structure based on the proposed technology stack and estimates.

5.4.1 Team Skillset

The section defines skills required from the different team member roles to implement the solution.

(describe team roles and required skillsets for these roles)

Role	Skillset
Business Analyst	• Business domain 1 • Business domain 2
Backend Developer	• Programming Language 1 • Technology 2 • Framework 3 • Tool 4 • Standard 5
UI Developer	• Programming Language 1 • Technology 2 • Framework 3 • Tool 4 • Standard 5

Role	Skillset
DevOps Engineer	• Programming Language 1 • Technology 2 • Framework 3 • Tool 4 • Standard 5

5.4.2 Team Structure

The section proposes a team/sub-team structure and work allocation for the implementation phase.

(describe teams and their responsibilities)

5.4.2.1 Core Team

The core team is responsible for:

- Development of <Component 1>
- Development of <Component 2>
- Support of <Application 1>

Role	Responsibility	Count	FTE
Project Manager	• Project management (coordinate the team, status reporting, communications with the Client team) • Gap analysis	1	0.5
Business Analyst	• Business analysis • Business requirements specification • Software requirements specification	1	1.0
Solutions Architect	• System analysis and design • System requirements specification	1	0.2

Role	Responsibility	Count	FTE
Technical Leader	• Technical leadership and communication • SCRUM Master • Code Reviews • Backend implementation	1	1.0
Sr. Backend Developer	• Technical communication • Code reviews • Backend implementation	2	1.0
Jr. Backend Developer	• Backend implementation	3	1.0
Int. UI Developer	• Web frontend design and implementation • UX prototyping	3	1.0

5.4.2.2 DevOps Team

The DevOps team is responsible for:

- Configuration of <Environment 1>
- Configuration of <Environment 2>
- Deployment and runtime monitoring of <Application 1>

Role	Responsibility	Count	FTE
DevOps Architect	• Operations analysis and design	1	0.2
Int. DevOps Engineer	• System deployment • Environment setup • Production operations support	3	1.0